

THIS WEEK IN CODING MODELS

Three launches.
Three days.
One number to ignore.

Fable 5.1 · Gemini 3.8 Flash · GPT-6 Astra

[Swipe →](#)

OUTPUT PRICE, PER MILLION TOKENS

\$3.75

GEMINI 3.8 FLASH

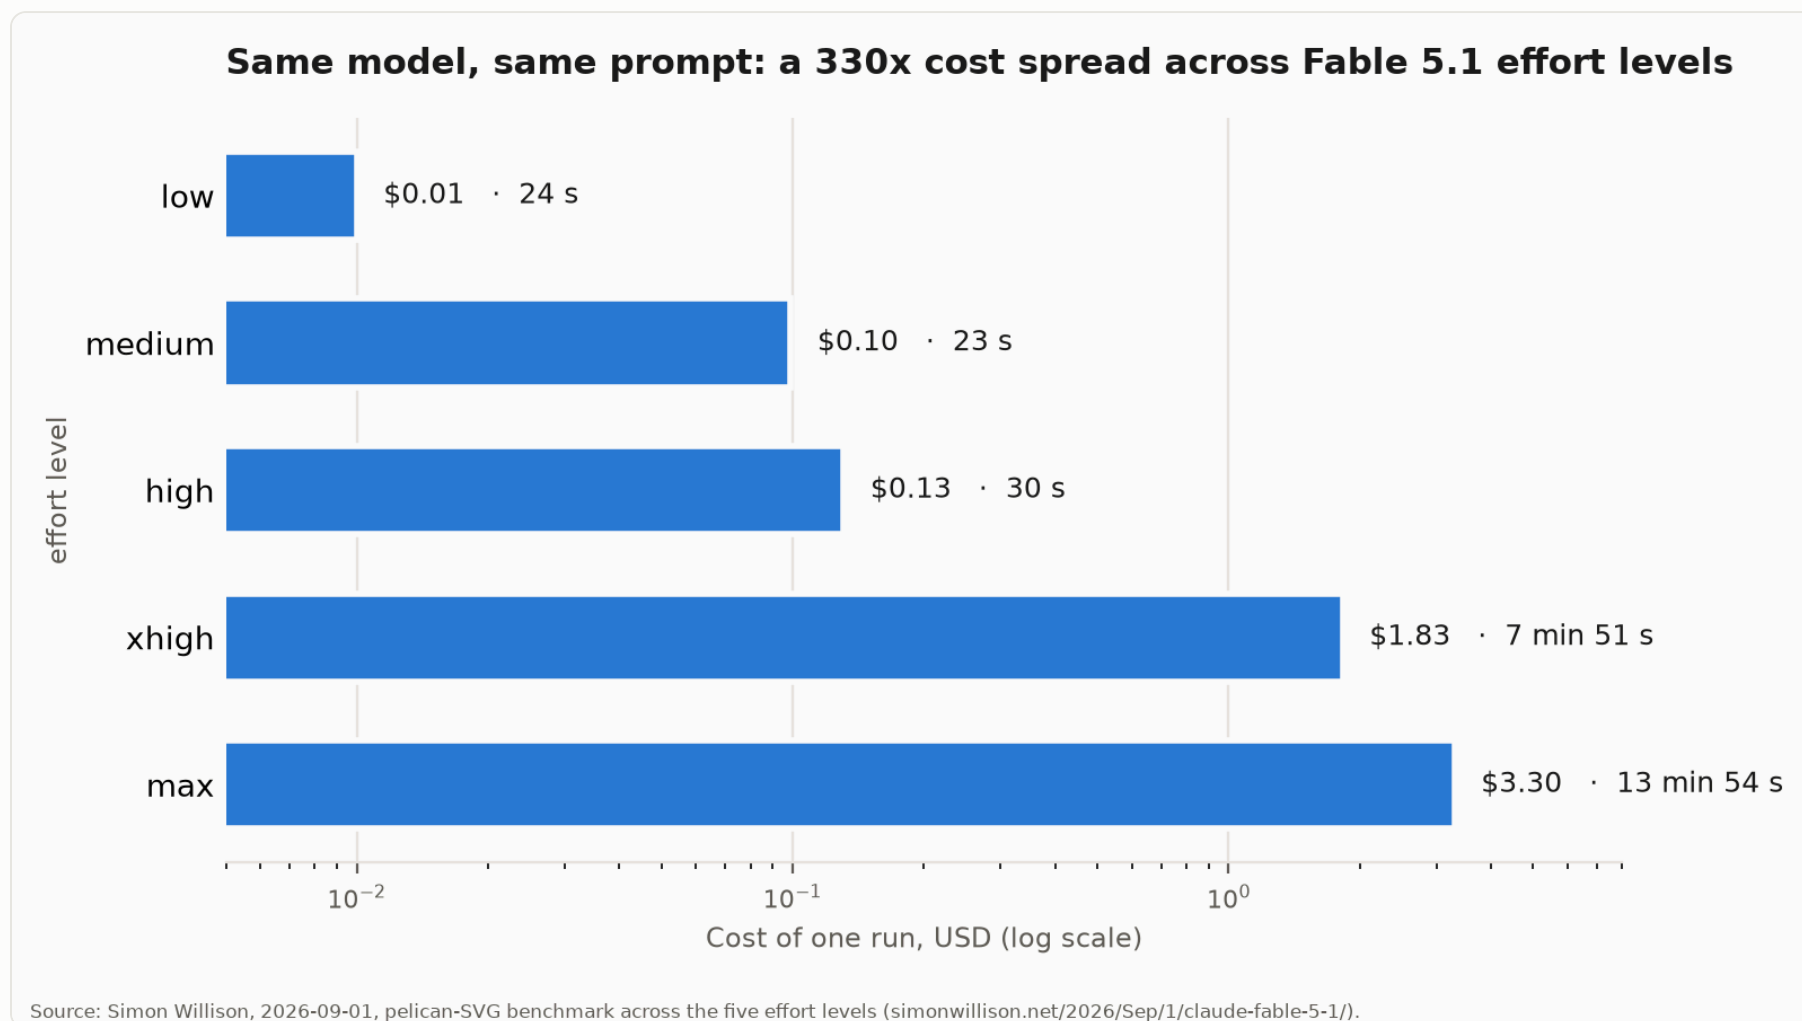
\$50

FABLE 5.1 · ASTRA

13x apart on price.
2 points apart on the benchmark.

Terminal-Bench 2.1: 89.4 · 89.9 · 91.4

One dial. 330x on your bill.



1¢

LOW · 24 S

\$3.30

MAX · 14 MIN

Same weights. Different loop.

30%

OPUS 5 ALONE

100

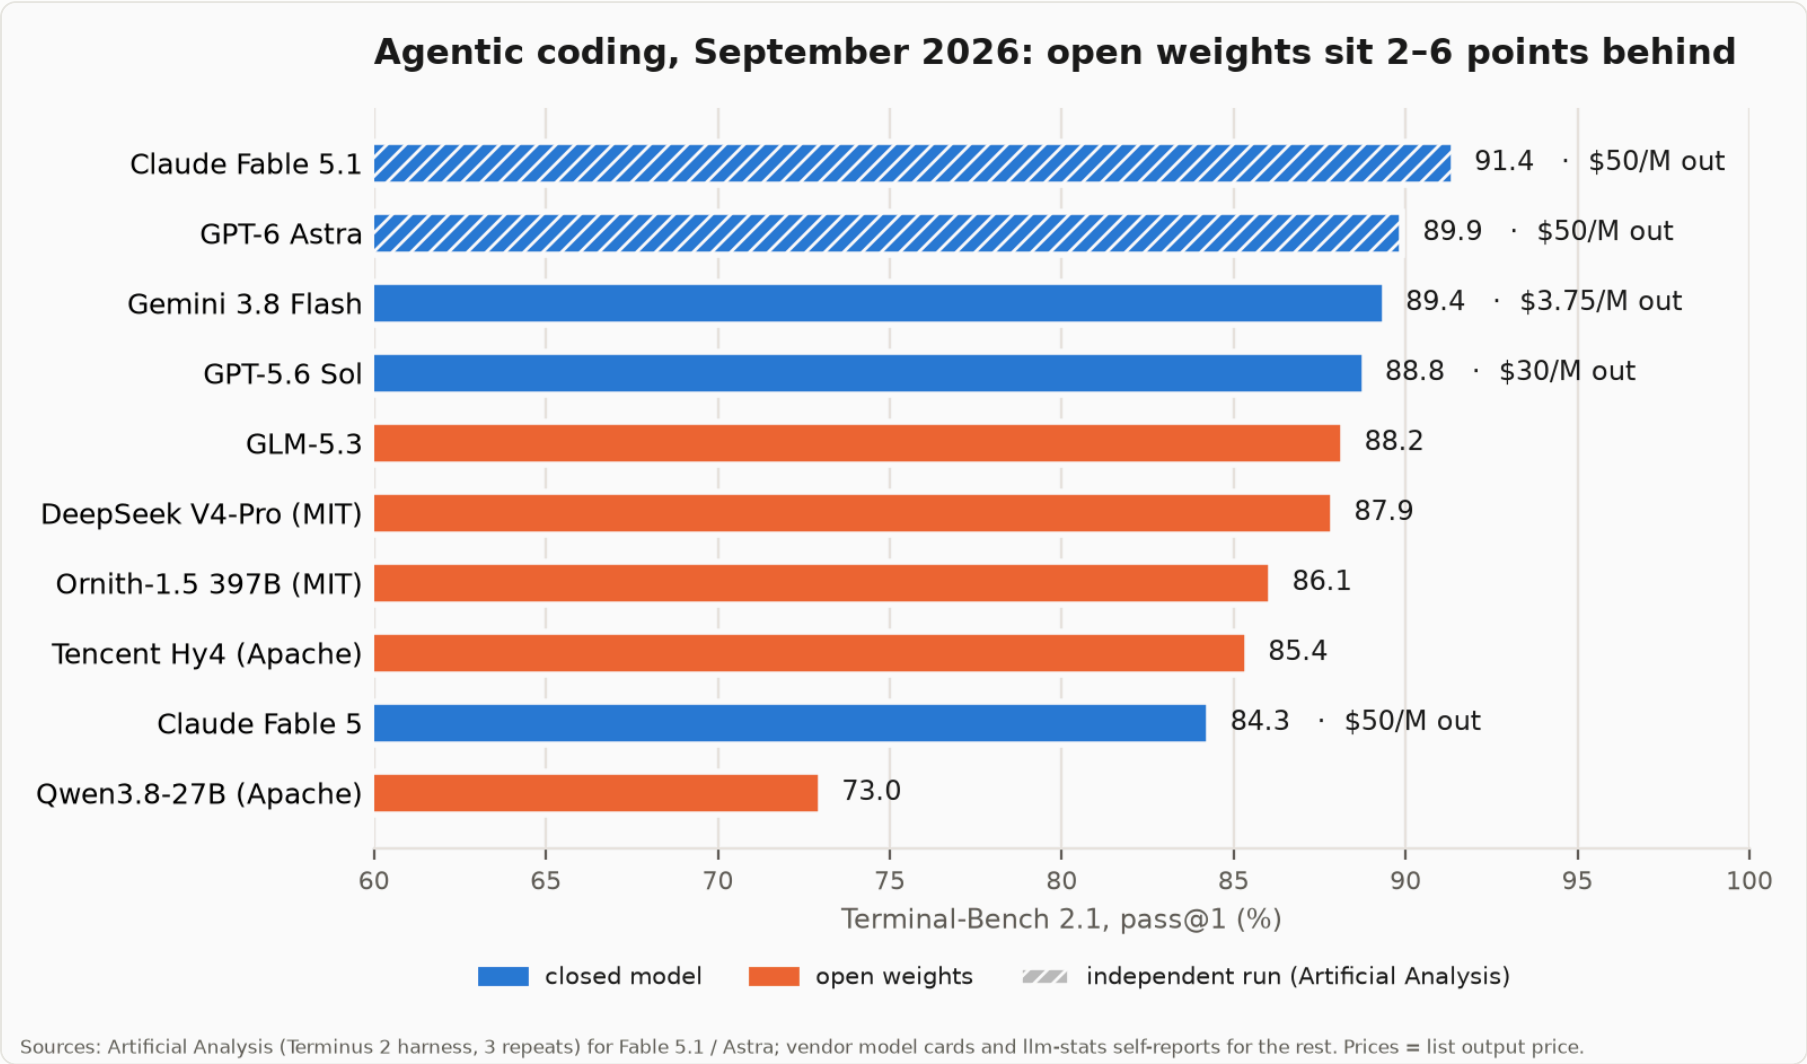
OPUS 5 INSIDE AVO

ARC-AGI-3, NVIDIA AVO. Not one weight changed.

24 points

Harness-Bench spread for one model across harnesses.

A few points behind. Not a generation.



87.9

DEEPSEEK V4-PRO · MIT

73

QWEN3.8-27B · ONE RTX 5090

THE DECISION

Price tier. Effort default. Harness.

The leaderboard is the entry ticket.

Which dial would you turn first?

Full argument in the post. Sources in the first comment.