

THIS WEEK IN CODING MODELS

Three coding model launches in three days. The leaderboard rank was the least useful number in every one.

Fable 5.1, Gemini 3.8 Flash, GPT-6 Astra. What actually touches your bill if you ship code with an agent.

[Swipe →](#)

LIST PRICE, PER MILLION TOKENS

Same job. 13x apart on price.

Sept 1 **Claude Fable 5.1** **\$10 / \$50**

cache reads down to \$0.25, the number that matters when your agent re-reads the repo for hours

Sept 2 **Gemini 3.8 Flash** **\$0.75 / \$3.75**


intro price, runs out on Dec 31

Sept 3 **GPT-6 Astra** **\$10 / \$50**

On Terminal-Bench 2.1 they sit within two points of each other: 91.4 and 89.9 in Artificial Analysis' own runs, 89.4 self-reported for Flash.

So the rank tells you almost nothing. Three other dials do.

1 cent at low. \$3.30 at max. Same model, same prompt.



Simon Willison ran one prompt through Fable 5.1 at every effort level. And Astra at its lowest setting beat every GPT-5.6 Sol setting on his test for under 10 cents.

Your default effort level is now a line item.

Same weights. Different loop around them.

30%

OPUS 5 ALONE

100


OPUS 5 INSIDE AVO

Claude Opus 5 on ARC-AGI-3, before and after NVIDIA wrapped it in the AVO agent architecture. Not one weight changed.

24 points

The spread Harness-Bench shows for a single model, depending on who wrote the loop around it.

A few points, not a generation.



DeepSeek V4-Pro (MIT) 87.9, Ornith-1.5 86.1, Tencent Hy4 85.4 on the same Terminal-Bench. Vendor numbers, sure.

And a 27B Qwen3.8 does 73 at full context on one RTX 5090 under your desk.

THE DECISION

Price tier. Effort default. Harness.

Choosing a coding model in September 2026 means choosing those three. The leaderboard is the entry ticket. The decision happens on the other three.

Which dial would you turn first: effort, harness, or going open?

Sources in the first comment.